

# NVIDIA Corporation

NVDA | October 4, 2026

## REDUCE

|                                       |                 |
|---------------------------------------|-----------------|
| Fair value (12-mo horizon reference)  | <b>\$186.00</b> |
| Current price (October 2, 2026 close) | <b>\$233.95</b> |
| Implied upside / downside             | <b>-21%</b>     |
| Verdict                               | <b>REDUCE</b>   |

## Investment Thesis

NVIDIA is the defining company of the AI era, and the numbers are almost difficult to believe: second-quarter fiscal 2027 revenue of \$96.2 billion, up 106% year over year; data-center revenue the overwhelming majority of the total; gross margins at 73.4%; Blackwell systems shipping at a pace management calls the fastest production ramp in semiconductor history; a \$10 billion-plus investment web spanning OpenAI, Anthropic, and Intel; and sovereign AI programs across the U.K., Germany, South Korea, and the Gulf. The CUDA software moat - two decades of developer lock-in, libraries, and tooling - remains the deepest competitive advantage in technology. We do not question that NVIDIA is a great company. We question whether any company in history has been worth what the market currently says NVIDIA is worth.

At \$233.95, the stock prices a world in which AI capital expenditure grows more or less indefinitely, in which NVIDIA captures the dominant share of every layer of that spend - compute, networking (with data-center networking up 199% year over year), and increasingly the systems and cloud commitments themselves - and in which gross margins near 75% persist against the combined competitive response of the largest technology companies on earth. Each of those assumptions deserves scrutiny. Hyperscalers are designing their own silicon (Google's TPUs, Amazon's Trainium, Microsoft's Maia) precisely to reduce the NVIDIA tax; AMD and a funded startup ecosystem contest the merchant-silicon layer; China - historically a meaningful market - is constrained by export controls; and the \$50 billion of manufacturing commitments and \$26 billion of cloud commitments on the balance sheet show a company stretching its own capital to sustain the demand narrative. Capex supercycles in technology have a perfect historical record: every one of them has ended, and the companies priced as though they would not have been the worst investments of the subsequent decade.

Our judgment, on a probability-weighted scenario DCF - 10-year horizon, 10% base discount rate reflecting technology cyclicality and customer concentration despite the moat, bear +250bp, bull -150bp - is fair value of \$186, about 21% below the current price. The base case is not bearish on AI: it assumes data-center revenue keeps growing at very high rates for several more years before moderating, CUDA sustains pricing power, and NVIDIA remains the dominant AI compute supplier with margins settling in the 60s rather than the 70s. The

bear case - which sits below the price, as required - is the capex-digestion scenario the market refuses to price: hyperscaler spending pauses after the buildout, custom silicon takes 15-20% share of AI compute, China revenue goes to zero, gross margins compress toward 60%, and the multiple compresses from growth-premium to semiconductor-cyclical. We rate the stock REDUCE. The AI revolution is real; revolutions do not repeal the business cycle, and this price assumes they do.

## Business Overview

NVIDIA designs accelerated-computing platforms: data-center GPUs and AI systems (the Blackwell and forthcoming Rubin architectures), networking (NVLink, InfiniBand, Spectrum-X Ethernet), the CUDA software stack, plus gaming, professional visualization, and automotive/robotics platforms. Data center is now roughly 90% of revenue - the company has, for economic purposes, become a picks-and-shovels monopoly on AI training and inference compute. Fiscal 2026 full-year revenue was \$215.9 billion, up 65%; the first two quarters of fiscal 2027 have already delivered \$177.8 billion combined, putting the company on a run-rate that would have been fantastical three years ago.

The moat has three layers. First, CUDA: twenty years of software investment that makes NVIDIA's hardware the default target for AI developers, with switching costs measured in rewritten codebases and retrained teams. Second, systems: the shift from selling chips to selling rack-scale systems (GB200/G300 NVL72) with integrated networking raised the competitive bar from silicon design to full-stack systems engineering, where NVIDIA's lead is widest. Third, cadence: annual architecture releases (Blackwell to Rubin) keep competitors perpetually a generation behind. Gross margins above 73% and operating cash flow of \$66.5 billion in nine months of fiscal 2026 are the financial expression of this dominance, funding \$36.7 billion of buybacks and a fortress balance sheet with \$60+ billion of liquidity against \$8.5 billion of debt.

The demand base is concentrated by any honest measure: a handful of hyperscalers (Microsoft, Google, Amazon, Oracle, Meta), xAI, sovereign buyers, and a growing enterprise tier account for the great majority of data-center revenue. Management discloses \$500 billion-plus of AI chip bookings visibility through 2026 and describes cloud GPU supply as 'sold out.' The strategic investments - up to \$10 billion in Anthropic, the OpenAI 10 GW partnership, \$5 billion in Intel, sovereign programs - serve the dual purpose of expanding the ecosystem and, inevitably, of supporting the demand narrative on which the valuation depends.

## Forward Outlook

Our forward view separates the AI thesis, which we believe, from the NVIDIA-stock thesis at this price, which we do not. AI compute demand will grow enormously through the decade: training runs keep scaling, inference demand compounds with deployment, and sovereign programs add a durable new buyer class. NVIDIA will remain the largest beneficiary of that growth - CUDA's developer lock-in does not dissolve, the systems lead is real, and Rubin extends the roadmap. Our base case has data-center revenue roughly tripling from fiscal 2026 levels over the explicit period before growth moderates to single digits. That is a bullish forecast by any historical standard for a company already at \$300 billion of run-rate revenue, and it still produces fair value 21% below the price.

The risks to the demand trajectory are not about whether AI matters but about the shape of the spending curve. First, capex digestion: hyperscalers are simultaneously building for uncertain future demand, and the history of telecom, dot-com, and memory capex cycles says the pause, when it comes, is abrupt - 20-30% order cuts in a single year are the norm, not the exception. Second, substitution: every hyperscaler has an economic incentive to move inference - the larger long-term workload - to custom silicon, and Google's TPUs already demonstrate

the model works at scale. Third, China: export controls have already removed a high-margin market, and further tightening is a one-way ratchet. Fourth, margins: 73%+ gross margins on merchant silicon are an invitation to competition that the entire industry is now RSVP-ing to, and systems-level competition (where NVIDIA bundles networking and software) only partly insulates the silicon margin.

Our judgment on timing: the current fiscal year and next are likely to remain very strong - backlog, bookings, and sold-out supply see to that - which is why we rate the stock REDUCE rather than SELL. The asymmetry is in the out-years: the market prices a smooth, decade-long capex ramp, while the historical base rate for technology capex cycles says the smooth ramp is the least likely path. When the digestion comes - and 'when,' not 'if,' is the honest framing - the combination of decelerating growth and multiple compression is what produces 30-50% drawdowns in prior-cycle leaders. We would rather be early in recognizing the asymmetry than precise in timing the turn.

## Valuation

We value NVIDIA on a probability-weighted scenario DCF over an explicit 10-year horizon, weighted 25% / 50% / 25%. The base discount rate is 10% - our standard rate, reflecting that even the deepest moat in technology does not repeal cyclicalities, customer concentration (a handful of hyperscalers drive the majority of data-center revenue), geopolitical exposure, and the inherent volatility of capex-driven demand. The bear case adds 250bp, the bull case subtracts 150bp with an 8% floor. Terminal growth is capped at 2.5% applied to normalized mid-cycle margins - emphatically not today's 73%+ gross margins, which reflect peak-cycle pricing power - and terminal value exceeding 70% of enterprise value is haircut and disclosed.

The bear case, sitting below the current price, is the capex-digestion scenario: hyperscaler spending pauses for 18-24 months after the current buildout wave, custom silicon captures 15-20% of AI compute (concentrated in inference), China revenue goes to zero under tighter controls, gross margins compress from the 70s toward 60% as competition intensifies, revenue actually declines in the worst years of the explicit period, and the terminal multiple reflects a mature semiconductor cyclical rather than a growth franchise. The base case assumes AI compute demand grows enormously but with cyclical air pockets: data-center revenue roughly triples over the decade, CUDA sustains a pricing premium, margins settle in the 60s, and NVIDIA exits the period as the dominant but no longer monopolistic AI compute supplier. The bull case assumes the market's implicit forecast - uninterrupted 20%+ growth deep into the 2030s, sustained 70%+ gross margins, successful extension of the franchise into sovereign and enterprise AI at current economics.

The probability-weighted fair value is \$186.00, our target, against the October 2, 2026 close of \$233.95 - about 21% downside. The valuation's message is the oldest one in growth investing: the business can be one of the great franchises of the era and the stock can still be 20% overvalued, because franchises are priced on fundamentals and this price is set by extrapolation.

## Key Risks

- Capex digestion: hyperscaler AI spending is cyclical; a pause after the current buildout wave would cut orders abruptly, and NVIDIA's revenue concentration makes the impact severe.
- Custom-silicon substitution: hyperscaler in-house chips (TPUs, Trainium, Maia) target exactly the high-volume inference workloads that underpin long-term demand.
- Customer concentration: a handful of hyperscalers and AI labs drive the majority of data-center revenue; capex decisions by 3-4 companies are the demand curve.

- China/export controls: further tightening would permanently remove a high-margin market; the ratchet moves one way.
- Margin compression: 73%+ gross margins attract the combined competitive response of the industry's largest R&D; budgets; AMD, startups, and custom silicon all attack the margin stack.
- Multiple compression: the stock embeds a growth-franchise multiple; any sustained deceleration reprices it toward semiconductor-cyclical multiples - a double hit with earnings.
- Balance-sheet stretch: \$50B+ of manufacturing commitments and large strategic investments show capital intensity rising with the narrative.

## What Would Change Our Mind

- Evidence that AI capex is through-cycle rather than cyclical - hyperscalers sustaining 20%+ capex growth through an actual economic downturn - would retire our central cyclicity objection.
- Custom-silicon efforts stalling or being abandoned by major hyperscalers would extend the duration of NVIDIA's pricing power materially.
- A 25-30% drawdown on capex-digestion fears, with CUDA's moat and the roadmap intact, would make the risk-reward positive and move us to BUY - prior-cycle leaders should be bought in the digestion, not the euphoria.
- Sustained 70%+ gross margins with 25%+ growth through a full semiconductor downturn would prove the business has structurally de-cyclicalized.
- Conversely, hyperscaler capex guidance cuts, lengthening GPU lead times normalizing to availability, or book-to-bill falling below 1 would confirm the cycle is turning and deepen the REDUCE thesis.

## Valuation Approach

Our valuation is a probability-weighted scenario DCF. We model bear, base, and bull cases over an explicit 10-year forecast horizon, weighted 25% / 50% / 25%. Discount rates are scenario-specific: the base rate reflects fundamental business risk (9% for stable franchises, 10% standard, 12% or higher for speculative situations); the bear case adds 250bp, the bull case subtracts 150bp (floor 8%). Terminal value assumes no more than 2.5% perpetual growth applied to normalized mid-cycle margins - never peak margins - and any terminal value exceeding 70% of enterprise value is haircut and disclosed. Bear cases are required to be genuinely adverse and to sit below the current price. Management guidance is never accepted at face value; it is independently tested and haircut where evidence warrants.

## Disclaimer

This report is independent research analysis for informational purposes only. It is not investment advice, a recommendation to buy or sell any security, or an offer to buy or sell any security. Valuations are inherently uncertain and scenario-dependent.